

Unit 1

Simple Linear Regression

Prof. Phil Schniter



THE OHIO STATE UNIVERSITY

ECE 5307: Introduction to Machine Learning, Sp23

Learning objectives

- Understand how to load datasets in Python using pandas
- Visualize data using a scatter plot
- Understand sample mean, variance, and covariance
- Describe a linear model for data
 - Identify the target variable and predictor
- Compute the least-squares fit using the regression formula
- Compute the R^2 measure-of-fit
- Visually assess goodness-of-fit and identify causes of poor fit

Outline

- Motivating Example: Predicting Automobile MPG
- Linear Model
- Least-Squares Fit: Problem Definition
- Sample Mean, Variance, Standard Deviation, Covariance
- Least-Squares Fit: The Solution
- Assessing Goodness-of-Fit

Example: Predicting automobile mpg

- We will consider the task of predicting the fuel efficiency (in mpg) of a car from features like # cylinders, horsepower, weight, etc.
- Demo in Jupyter notebook: `demo01_auto_mpg.ipynb`
 - Learn from demo, and test your understanding in the lab
- Uses data from UCI library: <https://archive.ics.uci.edu/ml/>
- Data is loaded using Python's `pandas` library
 - Pandas routines have many options!
 - Learn from examples; Google is your friend!
 - Pandas `read_csv` command creates a `dataframe` object with 3 main components:
 - `df.values`: numerical values in `Numpy` format
 - `df.columns`: column labels
 - `df.index`: row labels

Result of pandas' read_csv

```
import pandas as pd
import numpy as np
```

```
In [3]: names = ['mpg', 'cylinders', 'displacement', 'horsepower',
                 'weight', 'acceleration', 'model year', 'origin', 'car name']
```


```
In [4]: df = pd.read_csv('https://archive.ics.uci.edu/ml/machine-learning-databases/'+
                        'auto-mpg/auto-mpg.data',
                        header=None, delim_whitespace=True, names=names, na_values='?')
df.head(6)
```

```
Out[4]:
```

	mpg	cylinders	displacement	horsepower	weight	acceleration	model year	origin	car name
0	18.0	8	307.0	130.0	3504.0	12.0	70	1	chevrolet chevelle malibu
1	15.0	8	350.0	165.0	3693.0	11.5	70	1	buick skylark 320
2	18.0	8	318.0	150.0	3436.0	11.0	70	1	plymouth satellite
3	16.0	8	304.0	150.0	3433.0	12.0	70	1	amc rebel sst
4	17.0	8	302.0	140.0	3449.0	10.5	70	1	ford torino
5	15.0	8	429.0	198.0	4341.0	10.0	70	1	ford galaxie 500

Visualizing the data

- When possible, visualize the data before working with it.
- Python has MATLAB-like plotting features in the `matplotlib` module.



Visualizing the Data

We load the `matplotlib` module to plot the data in MATLAB.

```
In [15]: import matplotlib
import matplotlib.pyplot as plt
%matplotlib inline
```

First, let's extract some data columns and create arrays:




```
In [16]: xstr = 'displacement'
x = np.array(df[xstr])
y = np.array(df['mpg'])
```

Now we can create a scatter plot:

```
In [17]: plt.plot(x,y,'o')
plt.xlabel(xstr)
plt.ylabel('mpg')
plt.grid(True)
```

Postulating a model

- What relationships do you see?
- Is there a mathematical model relating the variables?
- How well can you predict mpg from these variables?




Outline

- Motivating Example: Predicting Automobile MPG
- **Linear Model**
- Least-Squares Fit: Problem Definition
- Sample Mean, Variance, Standard Deviation, Covariance
- Least-Squares Fit: The Solution
- Assessing Goodness-of-Fit

Describing and visualizing data

- y : the variable we are trying to predict
 - Many names: **target**, **label**, **regressand**, **response variable**, **dependent variable**
- x : the variable we are using for prediction
 - Many names: **feature**, **predictor**, **regressor**, **attribute**, **independent variable**
 - For now we consider a single variable. Next unit we'll consider multiple
- Data: the set of pairs $\{(x_i, y_i)\}_{i=1}^n$
 - Each pair is called a **sample**
 - We will use n for the number of samples
- A **scatter plot** is used to visualize the data pairs




Single-variable linear model

- Assume a linear relationship:

$$y \approx \beta_0 + \beta_1 x \triangleq \hat{y}$$

- β_1 : slope
- β_0 : intercept
- \hat{y} : prediction
- $\beta \triangleq [\beta_0, \beta_1]^T$ are the **parameters** of the model. Also called **coefficients** or **weights**
- Why this model?
 - easy to interpret & analyze
 - simple to optimize parameters (as we will see)
- When is this a good model?



The **regression line** $\hat{y}(x)$ is shown in red

Outline

- Motivating Example: Predicting Automobile MPG
- Linear Model
- **Least-Squares Fit: Problem Definition**
- Sample Mean, Variance, Standard Deviation, Covariance
- Least-Squares Fit: The Solution
- Assessing Goodness-of-Fit

Linear model: The residual

- Note: x does not *exactly* predict y :


$$y \approx \beta_0 + \beta_1 x$$

- Let's model this behavior explicitly using a **residual**, ϵ :

$$y = \beta_0 + \beta_1 x + \epsilon$$

- For the i th sample, we then have:

- **predicted value**: $\hat{y}_i = \beta_0 + \beta_1 x_i$
- **residual**: $\epsilon_i = y_i - \hat{y}_i$



The **residual** ϵ_i is the vertical deviation from y_i to the regression line

Least-squares fit

- How do we choose the model parameters $\beta = [\beta_0, \beta_1]^T$?
- Minimize the **residual sum of squares** (RSS):

$$\text{RSS}(\beta_0, \beta_1) \triangleq \sum_{i=1}^n \epsilon_i^2 = \sum_{i=1}^n (y_i - \hat{y}_i)^2$$

https://en.wikipedia.org/wiki/Residual_sum_of_squares

- Also called the **sum of squared errors** (SSE) and **sum of square residuals** (SSR)
- Note that ϵ_i and \hat{y}_i are implicitly functions of $[\beta_0, \beta_1]$
- The value of β that minimizes RSS is the **least-squares fit**.
 - Geometrically: minimizes sum of squared distances to the regression line.

The optimization approach to ML: A general recipe

General ML problem

- Assume a **model** with some **parameters**
- Get **training data**
- Choose a **loss function**
- Find parameters that **minimize** loss

Simple Linear Regression

- Linear model: $\hat{y} = \beta_0 + \beta_1 x$
- Data: $\{(x_i, y_i)\}_{i=1}^n$
- $\text{RSS}(\beta_0, \beta_1) \triangleq \sum_{i=1}^n (y_i - \hat{y}_i)^2$
- Find (β_0, β_1) that minimizes $\text{RSS}(\beta_0, \beta_1)$

Outline

- Motivating Example: Predicting Automobile MPG
- Linear Model
- Least-Squares Fit: Problem Definition
- Sample Mean, Variance, Standard Deviation, Covariance
- Least-Squares Fit: The Solution
- Assessing Goodness-of-Fit

Sample mean, variance, & standard deviation

If we summarize the data in the right way, we can easily design the optimal $\beta!$

■ Sample mean:

$$\bar{x} \triangleq \frac{1}{n} \sum_{i=1}^n x_i, \quad \bar{y} \triangleq \frac{1}{n} \sum_{i=1}^n y_i$$


■ Sample variance:

$$s_x^2 \triangleq \frac{1}{n} \sum_{i=1}^n (x_i - \bar{x})^2, \quad s_y^2 \triangleq \frac{1}{n} \sum_{i=1}^n (y_i - \bar{y})^2$$

- Note: some authors use $\frac{1}{n-1}$ to give an “unbiased estimate.” We’ll explain this later, in Unit 3.

■ Sample standard deviation (SD): s_x, s_y


- Simply the square-root of the sample variance



<https://en.wikipedia.org/wiki/Variance>

Visualizing sample mean & SD on scatter plot

- Sample means \bar{x} and \bar{y} :
 - The center of mass in each axis
- Standard deviations s_x and s_y :
 - The “spread” in each axis about the mean
 - If the data was Gaussian distributed...
 - 68% of points < 1 SD from mean
 - 95% of points < 2 SDs from mean
 - 99.7% of points < 3 SDs from mean
- What are your estimates of \bar{x} , \bar{y} , s_x , s_y from the above scatter plot?




Computing the sample mean & SD in Python

- We can exactly compute $\bar{x}, \bar{y}, s_x, s_y$ using the **Numpy** package in **Python**

```
In [27]: xm = np.mean(x)
          ym = np.mean(y)
          sxx = np.mean((x-xm)**2)
          syy = np.mean((y-ym)**2)
```

```
xm      = 104.47,      ym= 23.45
sqrt(sxx)= 38.44,    sqrt(syy)= 7.80
```



Sample covariance

- The **sample covariance** is


$$s_{xy} \triangleq \frac{1}{n} \sum_{i=1}^n (x_i - \bar{x})(y_i - \bar{y})$$

which indicates how “related” $\{x_i\}$ and $\{y_i\}$ are.

- The value s_{xy} is easier to interpret after normalization. This motivates the **sample (Pearson) correlation coefficient**:

$$\rho_{xy} \triangleq \frac{s_{xy}}{s_x s_y} \in [-1, 1]$$

- The property $\rho_{xy} \in [-1, 1]$ is a consequence of the Cauchy-Schwarz inequality.
- Example scatterplots for datasets with various ρ_{xy} :



https://en.wikipedia.org/wiki/Pearson_correlation_coefficient

Alternative expressions for sample variance & covariance

- Recall that the sample variance was defined as $s_x^2 \triangleq \frac{1}{n} \sum_{i=1}^n (x_i - \bar{x})^2$
- A very useful alternative formula can be found by expanding the square:

$$s_x^2 = \frac{1}{n} \sum_{i=1}^n x_i^2 - 2\bar{x} \frac{1}{n} \sum_{i=1}^n x_i + \bar{x}^2 = \frac{1}{n} \sum_{i=1}^n x_i^2 - \bar{x}^2 \quad \Rightarrow \quad \frac{1}{n} \sum_{i=1}^n x_i^2 = s_x^2 + \bar{x}^2$$

- Similarly, for the sample covariance we had $s_{xy} \triangleq \frac{1}{n} \sum_{i=1}^n (x_i - \bar{x})(y_i - \bar{y})$
- A useful alternative expression for that is

$$s_{xy} = \frac{1}{n} \sum_{i=1}^n x_i y_i - \bar{x} \bar{y} \quad \Rightarrow \quad \frac{1}{n} \sum_{i=1}^n x_i y_i = s_{xy} + \bar{x} \bar{y}$$

Notation

- We will use the following notation in this class (we'll try to be consistent)
- Note: some books/authors use different notations

Statistic	Notation	Formula	Python
sample mean	\bar{x}	$\frac{1}{n} \sum_{i=1}^n x_i$	xm
sample variance	$s_x^2 = s_{xx}$	$\frac{1}{n} \sum_{i=1}^n (x_i - \bar{x})^2$	sxx
sample standard deviation	$s_x = \sqrt{s_{xx}}$	$\sqrt{\frac{1}{n} \sum_{i=1}^n (x_i - \bar{x})^2}$	sx
sample covariance	s_{xy}	$\frac{1}{n} \sum_{i=1}^n (x_i - \bar{x})(y_i - \bar{y})$	sxy
sample correlation coefficient	ρ_{xy}	$\frac{s_{xy}}{s_x s_y}$	rhoxy

Outline

- Motivating Example: Predicting Automobile MPG
- Linear Model
- Least-Squares Fit: Problem Definition
- Sample Mean, Variance, Standard Deviation, Covariance
- Least-Squares Fit: The Solution
- Assessing Goodness-of-Fit

Minimizing RSS



- To minimize $\text{RSS}(\beta_0, \beta_1)$, we find the β_0 and β_1 that zero the **gradient**, i.e., the vector of partial derivatives:

$$\left[\frac{\partial \text{RSS}(\beta_0, \beta_1)}{\partial \beta_0}, \frac{\partial \text{RSS}(\beta_0, \beta_1)}{\partial \beta_1} \right]^T = \begin{bmatrix} 0 \\ 0 \end{bmatrix}$$

Because RSS is quadratic with a PSD Hessian, the zero-gradient point is guaranteed to be an RSS minimum (more on this later)

- After some manipulations (see next page), we obtain the optimal values

$$\beta_1 = \frac{s_{xy}}{s_{xx}} = \frac{\rho_{xy}s_y}{s_x}, \quad \beta_0 = \bar{y} - \beta_1\bar{x}$$



Minimizing RSS: Derivation

The minimum RSS is achieved by values of (β_0, β_1) that zero the gradient, i.e.,

$$0 = \frac{\partial \text{RSS}(\beta_0, \beta_1)}{\partial \beta_0} = \frac{\partial}{\partial \beta_0} \sum_{i=1}^n (y_i - \beta_0 - \beta_1 x_i)^2 = -2 \sum_{i=1}^n (y_i - \beta_0 - \beta_1 x_i) \quad (1)$$

$$0 = \frac{\partial \text{RSS}(\beta_0, \beta_1)}{\partial \beta_1} = \frac{\partial}{\partial \beta_1} \sum_{i=1}^n (y_i - \beta_0 - \beta_1 x_i)^2 = -2 \sum_{i=1}^n x_i (y_i - \beta_0 - \beta_1 x_i) \quad (2)$$

Starting with (1), we can multiply both sides by $-\frac{1}{2n}$ to give

$$0 = \frac{1}{n} \sum_{i=1}^n (y_i - \beta_0 - \beta_1 x_i) = \underbrace{\frac{1}{n} \sum_{i=1}^n y_i}_{\bar{y}} - \beta_0 - \beta_1 \underbrace{\frac{1}{n} \sum_{i=1}^n x_i}_{\bar{x}} \Leftrightarrow \boxed{\beta_0 = \bar{y} - \beta_1 \bar{x}} \quad (3)$$

Doing the same with (2) gives

$$0 = \frac{1}{n} \sum_{i=1}^n x_i (y_i - \beta_0 - \beta_1 x_i) = \frac{1}{n} \sum_{i=1}^n x_i y_i - \bar{x} \beta_0 - \beta_1 \frac{1}{n} \sum_{i=1}^n x_i^2 \quad (4)$$

Minimizing RSS: Derivation (continued)

Plugging in (3) into (4) gives

$$0 = \underbrace{\frac{1}{n} \sum_{i=1}^n x_i y_i - \bar{x} \bar{y}}_{s_{xy}} - \beta_1 \underbrace{\left(\frac{1}{n} \sum_{i=1}^n x_i^2 - \bar{x}^2 \right)}_{s_{xx}} \Leftrightarrow \boxed{\beta_1 = \frac{s_{xy}}{s_{xx}}}$$

To find the minimum RSS, we plug the optimal value of β_0 into the RSS definition to get


$$\begin{aligned} \text{RSS}(\beta_0, \beta_1) &= \sum_i (y_i - \beta_0 - \beta_1 x_i)^2 = \sum_i \left((y_i - \bar{y}) - \beta_1 (x_i - \bar{x}) \right)^2 \\ &= \underbrace{\sum_i (y_i - \bar{y})^2}_{ns_{yy}} - 2\beta_1 \underbrace{\sum_i (y_i - \bar{y})(x_i - \bar{x})}_{ns_{xy}} + \beta_1^2 \underbrace{\sum_i (x_i - \bar{x})^2}_{ns_{xx}} \end{aligned} \quad (5)$$

and then we plug the optimal value of β_1 into (5) to get

$$\begin{aligned} \text{RSS}(\beta_0, \beta_1) &= n \left(s_{yy} - 2 \frac{s_{xy}^2}{s_{xx}} + \frac{s_{xy}^2}{s_{xx}} \right) = n \left(s_{yy} - \frac{s_{xy}^2}{s_{xx}} \right) \\ &= n \left(1 - \frac{s_{xy}^2}{s_{xx} s_{yy}} \right) s_{yy} = n(1 - \rho_{xy}^2) s_{yy} \end{aligned}$$

Automobile demo

- Applied to our Python demo...



```

xm = np.mean(x)
ym = np.mean(y)
sxx = np.mean((x-xm)**2)
syy = np.mean((y-ym)**2)
syx = np.mean((y-ym)*(x-xm))
beta1 = syx/sxx
beta0 = ym - beta1*xm
Rsqr = syx**2/sxx/syy

```

beta0=39.94, beta1=-0.16

Rsqr=0.61

$$\widehat{\text{mpg}} = \beta_0 + \beta_1 \times \text{horsepower}$$

- Another good simple-linear-regression demo can be found at <https://stattrek.com/regression/regression-example.aspx?Tutorial=AP>

Outline

- Motivating Example: Predicting Automobile MPG
- Linear Model
- Least-Squares Fit: Problem Definition
- Sample Mean, Variance, Standard Deviation, Covariance
- Least-Squares Fit: The Solution
- Assessing Goodness-of-Fit

R^2 Goodness-of-fit

- Question: How to judge whether a predictor is doing “well” on a dataset?
- Answer: Use a *normalized* version of RSS:
 - RSS includes contributions from n training samples.
By considering RSS / n , we remove the dependence on n .
 - RSS / n depends on s_y^2 , the variance-of- y (i.e., if s_y^2 doubles then RSS / n doubles).
By considering $\frac{\text{RSS} / n}{s_y^2}$, we remove the dependence on s_y^2 .

- More commonly, we report

$$1 - \frac{\text{RSS} / n}{s_y^2} \triangleq R^2,$$

known as the “coefficient of determination”


https://en.wikipedia.org/wiki/Coefficient_of_determination

- $R^2 = 1$ implies that the predictor is perfect (i.e., $\hat{y}_i = y_i$)
- $R^2 = 0$ implies the predictor is no better than the trivial one (i.e., $\hat{y}_i = \bar{y}$)
- $R^2 < 0$ implies worse than trivial!
- With RSS-minimizing β_0 and β_1 , we know

$$\text{RSS}(\beta_0, \beta_1) = n(1 - \rho_{xy}^2) s_y^2 \Rightarrow R^2 = \rho_{xy}^2$$

Visualizing ρ_{xy}^2 from data $\{(x_i, y_i)\}_{i=1}^n$

- If $\rho_{xy}^2 \approx 1$, then LS linear model gives a very good fit
- If $\rho_{xy}^2 \approx 0$, then LS linear model gives a very poor fit
- Since LS coef $\beta_1 = \frac{\rho_{xy}s_y}{s_x}$, we have that $\text{sgn}(\beta_1) = \text{sgn}(\rho_{xy})$



← data with varying ρ_{xy} (and fixed $s_x = s_y$)



← data with $\rho_{xy} \in \{-1, 1\}$ (but varying s_y)

← data with $\rho_{xy} = 0$: linear prediction doesn't work

A better model for the automobile example?

- What if we predicted the inverse:

$$\frac{1}{\text{mpg}} \approx \beta_0 + \beta_1 \times \text{horsepower}$$
- This uses a *nonlinear data transformation* followed by linear regression
- Will explore this idea later in the course



Learning objectives

- Understand how to load datasets in Python using pandas
- Visualize data using a scatter plot
- Understand sample mean, variance, and covariance
- Describe a linear model for data
 - Identify the target variable and predictor
- Compute the least-squares fit using the regression formula
- Compute the R^2 measure-of-fit
- Visually assess goodness-of-fit and identify causes of poor fit