

Unit 3

Model-Order Selection and the Bias-Variance Tradeoff

Prof. Phil Schniter



THE OHIO STATE UNIVERSITY

ECE 5307: Introduction to Machine Learning, Sp23

Learning objectives

- Understand the problem of **model-order selection**
- Visually identify **underfitting** and **overfitting** in a scatterplot
- Understand the need to partition data into **training** and **testing** subsets
- Understand the **K-fold cross-validation** process
 - Use it to assess the test error for a given model
 - Use it to select the model order
- Understand the concepts of **bias**, **variance**, and **irreducible error**
 - Know how to compute each from synthetically generated data
 - Understand the **bias-variance tradeoff**

Outline


- Motivating Example: Polynomial Degree Selection
- Cross-validation
- The Bias-Variance Tradeoff
- From Model-Order Selection to Feature Selection

Polynomial regression

- Recall **polynomial regression** from last lecture
- Given data $\{(x_i, y_i)\}_{i=1}^n$, model target y as

$$y \approx \beta_0 + \beta_1 x + \cdots \beta_d x^d$$

- model parameters are $\beta = [\beta_0, \beta_1, \dots, \beta_d]^T$
- d is the **degree** of the polynomial
- given d , we can fit β using least-squares (multiple linear regression with $x_j \triangleq x^j$)
- Question: Can we select d from the data?
 - An example of “**model-order selection**”



Example with synthetic data


- We consider **synthetic data** generated using a noisy polynomial model
- $\{x_i\}$: 40 samples uniformly distributed in interval $[-1, 1]$
- $\{y_i\}$: generated as $y_i = f(x_i) + \epsilon_i$
 - $f(x) = \beta_0 + \beta_1x + \dots + \beta_dx^d$ with $d = 3$ for some “true” coefficients $\{\beta_j\}_{j=0}^d$
 - noise $\epsilon_i \sim \mathcal{N}(0, \sigma^2)$, independent over i
- Synthetic data is useful for analysis and experimentation!
 - We know the “ground truth”
 - Thus we can measure the performance of various predictors

```
# Import useful polynomial library
import numpy.polynomial.polynomial as poly

# True model parameters
beta = np.array([1,0.5,0,2]) # coefficients
wstd = 0.4 # noise std
dtrue = len(beta)-1 # true poly degree

# Independent data
nsamp = 100
xdat = np.random.uniform(-1,1,nsamp)

# Polynomial plus noise
y0 = poly.polyval(xdat,beta)
ydat = y0 + np.random.normal(0,wstd,nsamp)
```




Fitting with the true model order

- Could implement via `linear_model.LinearRegression` by constructing features $x_{ij} = x_i^j$ for $j = 1 \dots d$ and $i = 1 \dots n$
- Shortcut: `numpy.polynomial` package
- In any case, we need to choose d , the polynomial order for our model
- First, let's see what happens if $d = 3$, the true polynomial order
 - The LS fit looks very good!

```
d = 3
beta_hat = poly.polyfit(xdat,ydat,d)


# Plot true and prediction model
xp = np.linspace(-1,1,100)
yp = poly.polyval(xp,beta_hat)
yp_hat = poly.polyval(xp,beta_hat)
plt.xlim(-1,1)
plt.ylim(-1,3)
plt.plot(xp,yp,'r-',linewidth=2,label='True (dtrue=3)')
plt.plot(xp,yp_hat,'g-',linewidth=2,label='Predicted (d=3)')

# Plot data
plt.scatter(xdat,ydat,label='Data')
plt.legend(loc='upper left')
plt.grid()
plt.xlabel('x')
plt.ylabel('y')
```




Fitting with the wrong model order

$d = 1$: “underfitting”




$d = 15$: “overfitting”




another
illustration:


<https://medium.com/greyatom>



Underfitted



Good Fit/Robust



Overfitted


Is there a way to estimate the true d from the data $\{(x_i, y_i)\}_{i=1}^n$?

Select the model-order that minimizes training MSE?

Simple idea:

- For each hypothesized model order d :
 - Compute LS coefficients $\hat{\beta} \in \mathbb{R}^{d+1}$
 - Compute $\text{MSE}_{\text{train}}(d) \triangleq \frac{1}{n} \|\mathbf{y} - \mathbf{A}\hat{\beta}\|^2$
- Finally, pick the d that minimizes $\text{MSE}_{\text{train}}$

- Does this work?
 - $\text{MSE}_{\text{train}}(d)$ decreases with d
 - Suggests to choose d as large as possible
 - Leads to “overfitting”
- Why does this happen?



Overfitting

This is why we can't use training MSE to select the model order:

- Notice that we are choosing among a **nested set of models**

$$\beta = [\beta_0, \beta_1, 0, 0, 0, \dots]^T \quad \text{when } d = 1$$

$$\beta = [\beta_0, \beta_1, \beta_2, 0, 0, \dots]^T \quad \text{when } d = 2$$

$$\beta = [\beta_0, \beta_1, \beta_2, \beta_3, 0, \dots]^T \quad \text{when } d = 3$$

$$\vdots$$

$$\vdots$$


so that **each model is more capable than the previous model**

- The LS *training MSE gets no worse* as the model becomes more capable
 - Minimizing the training MSE leads to choosing the largest d

- When $d \geq n-1$, the least-squares $\text{MSE}_{\text{train}}(d) = 0$

$$\frac{1}{n} \|\mathbf{y} - \mathbf{A}\hat{\beta}\|^2 = 0$$

- When d is large, \hat{y}_i **tries to fit the training noise ϵ_i**
 - This is the main characteristic of **overfitting!**



<https://en.wikipedia.org/wiki/Overfitting>

Outline

- Motivating Example: Polynomial Degree Selection
- Cross-validation
- The Bias-Variance Tradeoff
- From Model-Order Selection to Feature Selection

To prevent overfitting, use cross-validation

Main idea:

Evaluate performance on “test data” that is independent of the training data

- Simplest version: Partition total dataset into **two subsets**, know as “folds”:


- n_{train} training samples: $\{(\mathbf{x}_{\text{train},i}, y_{\text{train},i})\}_{i=1}^{n_{\text{train}}}$
- $n_{\text{test}} = n - n_{\text{train}}$ test samples: $\{(\mathbf{x}_{\text{test},i}, y_{\text{test},i})\}_{i=1}^{n_{\text{test}}}$

- Then, for each hypothesized model-order d :

- Compute LS coefficients $\hat{\beta}$ from training data
- Predict the test targets: $\hat{y}_{\text{test},i} = [1 \ \mathbf{x}_{\text{test},i}^T] \hat{\beta}$
- Compute

$$\text{MSE}_{\text{test}}(d) \triangleq \frac{1}{n_{\text{test}}} \sum_{i=1}^{n_{\text{test}}} (y_{\text{test},i} - \hat{y}_{\text{test},i})^2$$

- Finally, **choose the d that minimizes $\text{MSE}_{\text{test}}(d)$** ... better but not perfect




Suggests model-order $d = 6$

K -fold cross-validation

Previously we considered splitting into training & test. More sophisticated approaches:

- **K -fold cross validation (CV)**
 - Partition (shuffled!) data into K folds
 - Train using $K-1$ folds, test using 1 fold
 - Repeat for each of K possible test folds
 - Typically use $K = 5$ or 10
 - Expensive: requires K parameter fits
 - **Good approx of true performance!**
- **Leave-one-out cross validation (LOOCV)**
 - Extreme case where $K = n$
(each test fold contains 1 sample!)
 - Very expensive unless n is small



<https://medium.com/@sebastiannorena>

Implementing K -fold cross-validation with `sklearn`

- Nested for-loop approach to CV:
 - Loop over $k = 1, \dots, K$ folds
 - Loop over $d = 1, \dots, D$ model-orders
 - Compute test $\text{MSE}_{d,k}$ for each order d & fold k
 - Average test MSE across K folds to get $\overline{\text{MSE}}_d \triangleq \frac{1}{K} \sum_{k=1}^K \text{MSE}_{d,k}$
 - Choose d giving smallest $\overline{\text{MSE}}_d$
- Can use `sklearn`'s `KFold` method to generate index sets for the folds!

```
# Create a k-fold object
k = 10
kfo = sklearn.model_selection.KFold(n_splits=k, shuffle=True)

# Model orders to be tested
dtest = np.arange(0,10)
nd = len(dtest)

MSEts = np.zeros((nd,k))

# Loop over the folds
for isplit, Ind in enumerate(kfo.split(xdat)): # enumerate n

    # Get the training data in the split
    Itr, Its = Ind
    #kfo.split( ) produced Ind, which contains a pair of ind
    xtr = xdat[Itr]
    ytr = ydat[Itr]
    xts = xdat[Its]
    yts = ydat[Its]

    # Loop over the model order
    for it, d in enumerate(dtest):

        # Fit data on training data
        beta_hat = poly.polyfit(xtr,ytr,d)

        # Measure MSE on test data
        yhat = poly.polyval(xts,beta_hat)
        MSEts[it,isplit] = np.mean((yhat-yts)**2)
```

Confidence intervals for K -fold cross-validation

- Problem: $\overline{\text{MSE}}_d = \frac{1}{K} \sum_{k=1}^K \text{MSE}_{d,k}$ may inaccurately estimate true MSE_d when K is small
- Can compute confidence bounds on $\overline{\text{MSE}}_d$ using the so-called **standard error (SE)**:


$$\text{SE}_d \triangleq \frac{\text{std}(\text{MSE}_d)}{\sqrt{K}}, \text{ where}$$

$$\text{std}(\text{MSE}_d) = \sqrt{\frac{1}{K-1} \sum_{k=1}^K (\text{MSE}_{d,k} - \overline{\text{MSE}}_d)^2}$$

- Above, $\frac{1}{K-1}$ gives “unbiased” estimate of $\text{var}(\text{MSE}_d)$, implemented via `ddof=1` below



```
MSE_mean = np.mean(MSEts,axis=1) #note mean is taken over
MSE_se = np.std(MSEts,axis=1,ddof=1)/np.sqrt(k)
plt.errorbar(dtest, MSE_mean, yerr=MSE_se, fmt='--')
plt.ylim(0,1.5)
plt.xlabel('Model order')
plt.ylabel('mean Test MSE')
plt.grid()
```

line shows $\overline{\text{MSE}}_d$, error bars show SE_d :



The one-standard-error rule

- Previously, said to choose d minimizing $\overline{\text{MSE}}_d$
 - But this sometimes overfits true model-order!
- Better approach: **one-standard-error (OSE) rule**
 - Use *simplest* model giving $\overline{\text{MSE}}_d$ within one SE of minimum $\overline{\text{MSE}}$
- Detailed procedure:
 - Set $d_{\min} = \arg \min_d \overline{\text{MSE}}_d$
 - Set $\overline{\text{MSE}}_{\text{tgt}} = \overline{\text{MSE}}_{d_{\min}} + \text{SE}_{d_{\min}}$
 - Find smallest d such that $\overline{\text{MSE}}_d \leq \overline{\text{MSE}}_{\text{tgt}}$
- In example on right: $d_{\min} = 8$, but OSE selects $d = 3$, which is the true model-order



Training, test, and validation data

- Sometimes you will see three folds of data...
 - 1 **Training data**: Used to train the model
 - used during design
 - 2 **Test data**: Used to tune the model hyperparameters (e.g., model order)
 - used during design
 - 3 **Validation data**: Used to estimate model performance on unseen data
 - used only *after* your design is finalized
 - Withheld from the contestants in ML contests (e.g., [kaggle](#))
- In many cases, the definitions of “test” and “validation” are swapped!
 - Always make sure you know the intended meaning
 - In this course, we'll use the definitions above
 - In this unit, we will consider only training & test data

Outline

- Motivating Example: Polynomial Degree Selection
- Cross-validation
- The Bias-Variance Tradeoff
- From Model-Order Selection to Feature Selection

Statistical learning theory

- With degree- d polynomial regression, we saw that
 - choosing d too small causes underfitting
 - choosing d too large causes overfitting
 - d can be optimized by minimizing the sample MSE_{test} through cross-validation
- But this is just one special case of a more general concept:
 - models that are too simple cause underfitting
 - models that are too complex cause overfitting
 - **model complexity** can be optimized by minimizing the *statistical* mean-squared error, $\text{MSE}_{\hat{y}}$
- From a theoretical perspective, we will see that ...
 - analyzing $\text{MSE}_{\hat{y}}$ leads to the **bias-variance equation**
 - minimizing $\text{MSE}_{\hat{y}}$ involves a **tradeoff** between bias and variance

Random Variables

- The concept of **randomness** is useful when we want to describe what *might* happen, or what *tends* to happen, in machine-learning experiments
- A **random variable** (RV) can generate different possible values, each with a prescribed probability
 - The values generated by a RV are called **realizations**
 - There are two types of RVs: **discrete** and **continuous**
- A **discrete RV** “ A ” takes on values from a countable set $\{a^{(1)}, a^{(2)}, \dots, a^{(K)}\}$
 - It’s described by its **probability mass function** (pmf) $p_A = [p_{A,1}, \dots, p_{A,K}]^T$
 - Here, $p_{A,k} \triangleq \Pr\{A = a^{(k)}\}$, where $0 \leq p_{A,k} \leq 1$ and $\sum_{k=1}^K p_{A,k} = 1$
- A **continuous RV** “ A ” takes on an uncountable number of values from \mathbb{R}
 - It’s described by its **cumulative distribution function** (cdf) $P_A(a) \triangleq \Pr\{A \leq a\}$
 - Also described by its **probability density function** (pdf) $p_A(\cdot) \triangleq \frac{d}{da} P_A(\cdot)$
 - Note $\Pr\{A \leq a\} = \int_{-\infty}^a p_A(a') da'$, where $p_A(a) \geq 0$ and $\int_{-\infty}^{\infty} p_A(a') da' = 1$
- See the document: **A Primer on Probability and Expectation**

Expectation

- **Expectation** $\mathbb{E}\{\cdot\}$ is the statistical mean of a random variable
- Formally, for any function $f(\cdot)$ and random variable A ,

$$\mathbb{E}\{f(A)\} = \sum_{k=1}^K f(a^{(k)}) p_{A,k} \quad \text{for a discrete RV}$$

$$\mathbb{E}\{f(A)\} = \int_{-\infty}^{\infty} f(a) p_A(a) da \quad \text{for a continuous RV}$$

Vector-valued random variables (e.g., $\mathbf{a} \in \mathbb{R}^M$) can be handled similarly

- We will avoid formalities for now and focus on two key properties of $\mathbb{E}\{\cdot\}$:
For any functions $f(\cdot)$ & $g(\cdot)$ and random variables A & B :
 - $\mathbb{E}\{c + d f(A)\} = c + d \mathbb{E}\{f(A)\}$ for deterministic c and d (by **linearity**)
 - $\mathbb{E}\{f(A)g(B)\} = \mathbb{E}\{f(A)\} \mathbb{E}\{g(B)\}$ for **independent** A and B :

$$p_{A,B}(a, b) = p_A(a)p_B(b) \quad \forall a, b$$
- **Variance** is defined as $\text{var}\{A\} \triangleq \mathbb{E}\{(A - \mathbb{E}\{A\})^2\} = \mathbb{E}\{A^2\} - \mathbb{E}\{A\}^2$

Conditional expectation

- **Conditional expectation** $\mathbb{E}\{f(A) \mid B = b\}$ is the mean value of a function $f(\cdot)$ of random variable A given that some other random variable B equals b .
- Formally, for any function $f(\cdot)$ and random variables $A, B \in \mathbb{R}$,

$$\mathbb{E}\{f(A) \mid B=b\} = \int_{-\infty}^{\infty} f(a) p_{A|B}(a|b) da \quad \text{for conditional pdf } p_{A|B}(\cdot|\cdot)$$

- Often we will see $\mathbb{E}\{f(A) \mid B\}$. This is like $\mathbb{E}\{f(A) \mid B=b\}$ but with b replaced by random variable B . Thus, $\mathbb{E}\{f(A) \mid B\}$ is random
- We will frequently encounter the **law of total expectation**, which says that $\mathbb{E}\{A\} = \mathbb{E}\{ \mathbb{E}\{A \mid B\} \}$.
 - This is often used to compute expectation one variable at a time.
- See the document: [A Primer on Probability and Expectation](#)

Statistical model

Setup for our theoretical analysis...

- True model: $y = f(\mathbf{x}) + \epsilon$ with $\mathbb{E}\{\epsilon\} = 0$ and $\text{var}\{\epsilon\} = \sigma^2$
 - noise ϵ is random with mean zero and variance σ^2 , and independent over draws
 - feature vectors \mathbf{x} also random, independent over draws, and independent of ϵ
 - model holds for both training $\{(\mathbf{x}_i, y_i, \epsilon_i)\}_{i=1}^n$ and test $(\mathbf{x}, y, \epsilon)$ quantities, which are independent of each other
- Prediction model: $\hat{y} = \hat{f}(\mathbf{x}; \hat{\beta})$ for some \hat{f} and trained coefficients $\hat{\beta}$
 - $\hat{\beta}$ was designed from training data $\{(\mathbf{x}_i, y_i)\}_{i=1}^n$, and thus is random
 - $\{\mathbf{x}, \epsilon, \hat{\beta}\}$ are mutually independent
- **Mean-squared error** on \hat{y} for a given \mathbf{x} : $\text{MSE}_{\hat{y}}(\mathbf{x}) \triangleq \mathbb{E} \{ (y - \hat{y})^2 \mid \mathbf{x} \}$
 - this expectation averages over ϵ (in y) and $\hat{\beta}$ (in \hat{y}), but holds \mathbf{x} fixed

Writing MSE in terms of bias and variance

We now analyze the **statistical MSE** for fixed test features \mathbf{x} :




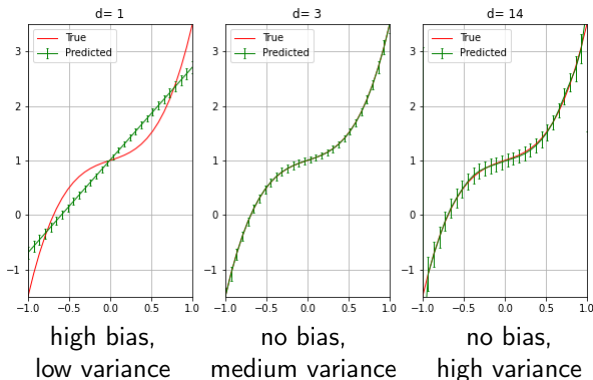
$$\begin{aligned}
 \text{MSE}_{\hat{y}}(\mathbf{x}) &\triangleq \mathbb{E} \{ (y - \hat{y})^2 \mid \mathbf{x} \} = \mathbb{E} \{ (\epsilon + f(\mathbf{x}) - \hat{f}(\mathbf{x}; \hat{\beta}))^2 \mid \mathbf{x} \} \\
 &= \mathbb{E} \{ \epsilon^2 + 2\epsilon(f(\mathbf{x}) - \hat{f}(\mathbf{x}; \hat{\beta})) + (f(\mathbf{x}) - \hat{f}(\mathbf{x}; \hat{\beta}))^2 \mid \mathbf{x} \} \quad \checkmark \text{ via linearity and independence of } \epsilon \text{ \& } \hat{\beta} \\
 &= \mathbb{E} \{ \epsilon^2 \mid \mathbf{x} \} + 2\mathbb{E} \{ \epsilon \mid \mathbf{x} \} \mathbb{E} \{ f(\mathbf{x}) - \hat{f}(\mathbf{x}; \hat{\beta}) \mid \mathbf{x} \} + \mathbb{E} \{ (f(\mathbf{x}) - \hat{f}(\mathbf{x}; \hat{\beta}))^2 \mid \mathbf{x} \} \\
 &= \sigma^2 + \mathbb{E} \{ (f(\mathbf{x}) - \mathbb{E}[\hat{f}(\mathbf{x}; \hat{\beta}) \mid \mathbf{x}] + \mathbb{E}[\hat{f}(\mathbf{x}; \hat{\beta}) \mid \mathbf{x}] - \hat{f}(\mathbf{x}; \hat{\beta}))^2 \mid \mathbf{x} \} \\
 &= \sigma^2 + \mathbb{E} \{ (f(\mathbf{x}) - \mathbb{E}[\hat{f}(\mathbf{x}; \hat{\beta}) \mid \mathbf{x}])^2 \mid \mathbf{x} \} + \mathbb{E} \{ (\mathbb{E}[\hat{f}(\mathbf{x}; \hat{\beta}) \mid \mathbf{x}] - \hat{f}(\mathbf{x}; \hat{\beta}))^2 \mid \mathbf{x} \} \\
 &\quad + 2(f(\mathbf{x}) - \mathbb{E}[\hat{f}(\mathbf{x}; \hat{\beta}) \mid \mathbf{x}]) \mathbb{E} \{ \mathbb{E}[\hat{f}(\mathbf{x}; \hat{\beta}) \mid \mathbf{x}] - \hat{f}(\mathbf{x}; \hat{\beta}) \mid \mathbf{x} \} \\
 &= \underbrace{\sigma^2}_{\text{irreducible error}} + \underbrace{(f(\mathbf{x}) - \mathbb{E}[\hat{f}(\mathbf{x}; \hat{\beta}) \mid \mathbf{x}])^2}_{\text{bias}_{\hat{y}}(\mathbf{x}) \triangleq \mathbb{E}\{y - \hat{y} \mid \mathbf{x}\}} + \underbrace{\mathbb{E} \{ (\hat{f}(\mathbf{x}; \hat{\beta}) - \mathbb{E}[\hat{f}(\mathbf{x}; \hat{\beta}) \mid \mathbf{x}])^2 \mid \mathbf{x} \}}_{\text{var}\{\hat{y} \mid \mathbf{x}\}}
 \end{aligned}$$

We can go one step further and take the mean of $\text{MSE}_{\hat{y}}(\mathbf{x})$ over random \mathbf{x} :

$$\text{MSE}_{\hat{y}} \triangleq \mathbb{E} \{ \text{MSE}_{\hat{y}}(\mathbf{x}) \} = \sigma^2 + \mathbb{E}\{\text{bias}_{\hat{y}}(\mathbf{x})^2\} + \mathbb{E}\{\text{var}\{\hat{y} \mid \mathbf{x}\}\}$$

A bias-variance experiment for polynomial models

- Polynomial demo
- Red curve:
 $f(x) = \mathbb{E}\{y|x\}$
- Solid green curves:
 $\mathbb{E}\{\hat{y} | x\}$
 - $\text{bias}_{\hat{y}}(x)$ is gap between red & green curves at x
- Green error-bars:
 $\sqrt{\text{var}\{\hat{y} | x\}}$
- The $\mathbb{E}\{.\}$ and $\text{var}\{.\}$ are approximated by sample-averaging 100 independent train/test experiments



A bias-variance analysis of LS linear regression

- Consider **noisy linear training data** $\{(\mathbf{x}_i, y_i)\}_{i=1}^n$:
 - $y_i = f(\mathbf{x}_i) + \epsilon_i$ with $f(\mathbf{x}_i) = \beta_0 + \sum_{j=1}^{d_{\text{true}}} \beta_j x_{ij}$ and $\epsilon_i \sim \mathcal{N}(0, \sigma^2)$
 and the d -term **linear regression model**:
 - $\hat{y} = \hat{f}(\mathbf{x}; \hat{\beta}) = \hat{\beta}_0 + \sum_{j=1}^d \hat{\beta}_j x_j$ with LS weights $\hat{\beta}$
- Result 1: If $n < d+1$, then $\hat{\beta}$ is not unique, so LS solution undefined
- Result 2: If $n \geq d+1$ and $d < d_{\text{true}}$, then \hat{y} will be **biased** due to **underfitting**
- Result 3: If $n \geq d+1$ and $d \geq d_{\text{true}}$, then \hat{y} is **unbiased**, i.e.,

$$\text{bias}_{\hat{y}}(\mathbf{x}) = \mathbb{E}\{\hat{y} - y \mid \mathbf{x}\} = 0$$

- Result 4: If $n \gg d$ and $d \geq d_{\text{true}}$ and \mathbf{x} has same distribution as $\{\mathbf{x}_i\}$,


$$\mathbb{E}\{\text{var}\{\hat{y} \mid \mathbf{x}\}\} \approx \frac{d+1}{n} \sigma^2 \quad \text{so} \quad \begin{cases} \text{variance increases linearly with \# model parameters} \\ \text{variance decreases inversely with \# training samples} \end{cases}$$

Details in handout “Bias and Variance Analysis of Multiple Linear Regression”

The bias-variance tradeoff for general regression

$$\text{MSE}_{\hat{y}} = \sigma^2 + \mathbb{E}\{\text{bias}_{\hat{y}}(\mathbf{x})^2\} + \mathbb{E}\{\text{var}\{\hat{y} | \mathbf{x}\}\}$$

- We saw two examples of how $\text{MSE}_{\hat{y}}$ changes with model complexity:
 - polynomials: polynomial degree d
 - linear regression: # features d
- Similar trends hold for general models!
 - There exists a **tradeoff** between bias and variance
- The **optimal** model complexity depends on
 - the true model complexity (which affects bias)
 - the number of training samples (which affects variance)



simpler models
less parameters
underfitting



richer models
more parameters
overfitting



Example: bias and variance of the sample estimators

- Consider random variable Z with mean $\mathbb{E}\{Z\} = \mu$ and variance $\text{var}\{Z\} = v$
- Say we want to estimate μ from n independent realizations $\{z_i\}_{i=1}^n$ of Z
 - It's common to use the sample mean $\bar{z} \triangleq \frac{1}{n} \sum_{i=1}^n z_i$
 - Can show that $\mathbb{E}\{\bar{z}\} = \mu$. Thus \bar{z} is an **unbiased** estimate of μ
 - Can also show that $\text{var}\{\bar{z}\} = \mathbb{E}\{(\bar{z} - \mu)^2\} = v/n$. **Standard error** = $\sqrt{\text{var}\{\bar{z}\}}$
- Now say we want to estimate v from independent realizations $\{z_i\}_{i=1}^n$ of Z
 - First consider sample variance $s_{zz} \triangleq \frac{1}{n} \sum_{i=1}^n (z_i - \bar{z})^2$
 - Can show that $\mathbb{E}\{s_{zz}\} = \frac{n-1}{n}v$. Thus s_{zz} is a **biased** estimate of v
 - But $\xi_{zz} \triangleq \frac{1}{n-1} \sum_{i=1}^n (z_i - \bar{z})^2$ is an **unbiased** estimate of v
 - Although s_{zz} is biased, it has a lower mean-squared error. If $Z \sim \mathcal{N}(\mu, v)$ then

$$\mathbb{E}\{(s_{zz} - v)^2\} = \frac{2n+1}{n^2}v^2 < \frac{2}{n-1}v^2 = \mathbb{E}\{(\xi_{zz} - v)^2\}$$

Details in handout “On the Bias and Variance of the Sample Estimators”

Outline

- Motivating Example: Polynomial Degree Selection
- Cross-validation
- The Bias-Variance Tradeoff
- From Model-Order Selection to Feature Selection

Feature selection: A generalization of model-order selection

- So far, we discussed model-*order* selection (e.g., polynomial degree d)
 - Select between several models, each with a different complexity
 - For example, $y \approx \beta_0 + \beta_1 x_1$
 $y \approx \beta_0 + \beta_1 x_1 + \beta_2 x_2$
 $y \approx \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \beta_3 x_3$
- More generally, given d total features $\{x_j\}_{j=1}^d$, we might wonder **which subset of features** works best for predicting y
 - Called “**feature selection**”
 - Given d features (plus intercept), there are 2^d possible subsets
- How do we choose the best subset?
 - Can use cross-validation to choose between models
 - but need to manage computational complexity...
- Discussed further in the next unit ...

Learning objectives

- Understand the problem of **model-order selection**
- Visually identify **underfitting** and **overfitting** in a scatterplot
- Understand the need to partition data into **training** and **testing** subsets
- Understand the **K-fold cross-validation** process
 - Use it to assess the test error for a given model
 - Use it to select model order
- Understand the concepts of **bias**, **variance**, and **irreducible error**
 - Know how to compute each from synthetically generated data
 - Understand the **bias-variance tradeoff**