

# Unit 11

## Principal Component Analysis

Prof. Phil Schniter



THE OHIO STATE UNIVERSITY

ECE 5307: Introduction to Machine Learning, Sp23

# Learning objectives

- Recognize need for feature dimensionality reduction
- Understand PCA as RSS-minimizing linear approximation
  - Understand orthogonal projection
  - Recognize PCA as subspace fitting
  - Understand the role of the data-covariance eigenvectors in PCA
  - Know how to measure PCA performance using PoV
- Understand how to compute PCA using the SVD
- Understand how the PCA coefficients can be used in supervised learning tasks
- Understand how PCA can be used for data visualization
- Be familiar with t-SNE and UNET, non-linear data-visualization techniques

# Outline

- Dimensionality Reduction
- Principal Component Analysis (PCA)
- Computing PCA via the SVD
- Python Example: Eigenfaces and PCA-based Classification
- Data Visualization using PCA, t-SNE, and UMAP

# Dimensionality reduction

- Many modern datasets have very high dimension  $d$
- We would like to reduce the dimension (if possible) ...
  - to simplify classification/regression tasks
  - to save memory/storage space
  - to help visualize structure in data
- In this unit, we focus on dimensionality reduction via PCA
  - PCA is RSS-optimal *linear* dimensionality reduction
- We also briefly describe t-SNE and UMAP
  - Nonlinear dimensionality reduction techniques often used for visualization

# Data representation

- Dataset:  $\{\mathbf{x}_i\}_{i=1}^n$ 
  - Each sample has  $d$  features:  $\mathbf{x}_i = [x_{i1}, \dots, x_{id}]^T \in \mathbb{R}^d$
  - Can represent dataset using the matrix  $\mathbf{X} = [\mathbf{x}_1, \dots, \mathbf{x}_n]^T \in \mathbb{R}^{n \times d}$
  - Will assume data is centered (i.e., mean was removed, so  $\sum_{i=1}^n \mathbf{x}_i = \mathbf{0}$ )
- Note: there are *no targets/labels* here!
  - Either they don't exist, or we are ignoring them
  - This is known as **unsupervised learning**
  - Until now, we've focused on **supervised learning**, e.g., classification, regression
- What if data dimension  $d$  is very large?
  - Can we reduce the dimension to ease further data processing?

## Example: Face data

- Face images can be high dimensional
  - We'll use the “**Labeled Faces in the Wild (LFW)**” dataset from 2007
  - These images contain  $d = 50 \times 37 = 1850$  pixels
  - Modern face datasets are much larger, e.g., up to 1 million pixels
- As we will see, face images can be well approximated using a few coefficients
  - Can be “compressed”!
- How exactly do we do this?



# Loading the data

- The LFW face dataset is built into **sklearn**
  - The full collection contains  $n = 13000$  images (from news stories in 2000s)
  - By requiring  $\geq 70$  faces per person, we extract a subset of 1288 images

```
from sklearn.datasets import fetch_lfw_people  
lfw_people = fetch_lfw_people(min_faces_per_person=70, resize=0.4)
```

Image size =  $50 \times 37 = 1850$  pixels

Number of samples = 1288

Number of classes = 7

- Some example faces:

Donald Rumsfeld



Colin Powell



George W Bush



Gerhard Schroeder



# Outline

- Dimensionality Reduction
- Principal Component Analysis (PCA)
- Computing PCA via the SVD
- Python Example: Eigenfaces and PCA-based Classification
- Data Visualization using PCA, t-SNE, and UMAP



# PCA — Main ideas

- Main idea 1: **Linearly approximate** each feature vector  $\mathbf{x}_i \in \mathbb{R}^d$  as follows:

$$\mathbf{x}_i \approx \mathbf{B}\mathbf{z}_i \text{ with } \mathbf{z}_i \in \mathbb{R}^R, \quad i = 1 \dots n$$

- The columns of  $\mathbf{B} \in \mathbb{R}^{d \times R}$  form a “dictionary” with  $R$  elements
- $\mathbf{z}_i$  contains  $R$  coefficients to linearly combine the dictionary elements for image  $i$
- $\mathbf{B}\mathbf{z}_i$  is a *linear* approximation of  $\mathbf{x}_i$ . (Linear is chosen for simplicity)
- $R$  is the “rank” of the approximation, where  $1 \leq R < d$  (and ideally  $R \ll d$ )
- Main idea 2: Design the approximation to **minimize RSS**:

$$(\hat{\mathbf{B}}, \{\hat{\mathbf{z}}_i\}_{i=1}^n) = \arg \min_{\mathbf{B}, \{\mathbf{z}_i\}} \left\{ \sum_{i=1}^n \|\mathbf{x}_i - \mathbf{B}\mathbf{z}_i\|^2 \right\}$$

- This is known as “**principal component analysis**” (PCA)
- RSS is used for simplicity
  - Caveat: RSS may be poorly matched to downstream processing (e.g., classification)

# PCA — Solution

- The optimal  $R$ -element approximation dictionary is



$$\boxed{\hat{B} = V_R} \triangleq [v_1, \dots, v_R]$$

where  $\{v_r\}_{r=1}^R$  are the eigenvectors corresponding to the  $R$  largest eigenvalues  $\{\lambda_r\}_{r=1}^R$  of the sample covariance matrix  $Q$ :

$$Q \triangleq \frac{1}{n} \sum_{i=1}^n x_i x_i^T = \sum_{j=1}^d \lambda_j v_j v_j^T \quad \text{since } \{x_i\} \text{ is centered}$$

- The optimal coefficients for the approximation of  $x_i$  are  $\boxed{\hat{z}_i = V_R^T x_i}$

- The PCA approximation projects  $x_i \in \mathbb{R}^d$  onto the subspace spanned by  $\{v_r\}_{r=1}^R$ , which are known as the  $R$  “principal components”



# PCA — Derivation ...

- Our derivation of PCA will proceed in two steps:
  - 1 Optimize the coefficients  $\{z_i\}$  for an arbitrary fixed dictionary  $B$
  - 2 Optimize the dictionary  $B$
- When optimizing  $z_i$ , we encounter the familiar LS problem from Unit 2:

$$\hat{z}_i = \arg \min_{z_i} \|x_i - Bz_i\|^2 = (B^T B)^{-1} B^T x_i$$

- Plugging  $\{\hat{z}_i\}_{i=1}^n$  back into the original problem yields

$$\begin{aligned} \hat{B} &= \arg \min_B \left\{ \sum_{i=1}^n \|x_i - B(B^T B)^{-1} B^T x_i\|^2 \right\} \\ &= \arg \min_B \left\{ \sum_{i=1}^n \underbrace{\| (I - B(B^T B)^{-1} B^T) x_i \|^2}_{\triangleq P_B^\perp} \right\} \end{aligned}$$

- As we show next, we can interpret  $P_B^\perp$  as an **orthogonal projection matrix**...

# Orthogonal projection

- Consider the **subspace**  $\text{colsp}(B) \triangleq \{Bz \text{ s.t. } z \in \mathbb{R}^R\} \subset \mathbb{R}^d$ 
  - the set of all linear combinations of the columns of  $B$

- The **orthogonal projection** of  $x \in \mathbb{R}^d$  onto  $\text{colsp}(B)$  ...

- is the closest vector to  $x$  within  $\text{colsp}(B)$ . So,

$$x = \hat{x} + e, \text{ where } \hat{x} \in \text{colsp}(B), e \perp \text{colsp}(B),$$

- which can be computed via


$$\hat{x} = P_B x \text{ and } e = P_B^\perp x$$

- using the orthogonal **projection matrices**

$$P_B \triangleq B(B^\top B)^{-1}B^\top \text{ and } P_B^\perp \triangleq I - P_B$$

- Such matrices are **symmetric** and **idempotent**, i.e.,  $P_B = P_B^\top$  &  $P_B = P_B^2$

- Note:  $P_B$  has  $R$  eigenvalues = 1, and all other eigenvals = 0, while  $P_B^\perp$  has  $R$  eigenvalues = 0, and all other eigenvals = 1



# The role of projection in PCA


- Back to the PCA problem:

$$\hat{B} = \arg \min_B \left\{ \sum_{i=1}^n \|P_B^\perp x_i\|^2 \right\}$$

- We now recognize  $P_B^\perp x_i$  as projection error
- Thus, PCA chooses  $B$  to minimize the sum-squared projection error
- To further understand  $\hat{B}$ , we reformulate the above cost:

$$\begin{aligned} J(B) &\triangleq \sum_i \|P_B^\perp x_i\|^2 = \sum_i x_i^\top (P_B^\perp)^\top P_B^\perp x_i = \sum_i x_i^\top P_B^\perp x_i = \sum_i \text{tr}(x_i^\top P_B^\perp x_i) \\ &= \sum_i \text{tr}(P_B^\perp x_i x_i^\top) = \text{tr}\left(P_B^\perp \sum_i x_i x_i^\top\right) = n \text{tr}\left(P_B^\perp \underbrace{\frac{1}{n} \sum_{i=1}^n x_i x_i^\top}_{\text{sample covariance mtx } Q}\right) \end{aligned}$$

where  $\text{tr}(A) \triangleq \sum_j [A]_{jj}$  and  $\text{tr}(AC) = \text{tr}(CA)$ .



# Eigen-decomposition of sample covariance matrices

- The sample covariance mtx  $Q$  is **positive semi-definite**, i.e.,  $x^T Q x \geq 0 \forall x$ 
  - Proof:  $x^T Q x = x^T (\frac{1}{n} \sum_i x_i x_i^T) x = \frac{1}{n} \sum_i (x^T x_i)(x_i^T x) = \frac{1}{n} \sum_i (x^T x_i)^2 \geq 0$

- All positive semi-definite matrices have an **eigen-decomposition** of the form

$$Q = V \Lambda V^T \quad \text{where} \quad \begin{cases} V \text{ is orthogonal (i.e., } VV^T = I_d = V^T V) \\ \Lambda = \text{Diag}(\lambda_1, \dots, \lambda_d) \text{ with } \lambda_j \geq 0 \forall j \end{cases}$$

Without loss of generality, we will assume  $\{\lambda_j\}$  are **sorted** from large to small

## Theorem (Eckart-Young, 1936)

The optimal  $B \in \mathbb{R}^{d \times R}$  is constructed from the  $R$  **principal eigenvectors** of  $Q$ :

$$\hat{B} = \arg \min_B n \text{tr}(P_B^\perp Q) = [v_1, \dots, v_R] \triangleq V_R,$$

More precisely, the optimal  $\hat{B}$  is *any*  $B$  for which  $\text{colsp}(B) = \text{colsp}(V_R)$

# A simple proof of Eckart-Young

- Recall that we want to minimize the RSS cost

$$J(\mathbf{B}) = n \operatorname{tr}(\mathbf{P}_B^\perp \mathbf{Q}) = n \operatorname{tr}((\mathbf{I}_d - \mathbf{P}_B)\mathbf{Q}) = n \operatorname{tr}(\mathbf{Q}) - n \operatorname{tr}(\mathbf{P}_B \mathbf{Q})$$

- Equivalently, we can *maximize* the utility

$$U(\mathbf{B}) \triangleq \operatorname{tr}(\mathbf{P}_B \mathbf{Q}) = \operatorname{tr}(\mathbf{P}_B \mathbf{V} \mathbf{\Lambda} \mathbf{V}^\top) = \operatorname{tr}(\mathbf{V}^\top \mathbf{P}_B \mathbf{V} \mathbf{\Lambda}) = \sum_{j=1}^d \alpha_j \lambda_j$$

$$\text{for } \alpha_j \triangleq [\mathbf{V}^\top \mathbf{P}_B \mathbf{V}]_{jj} = \mathbf{v}_j^\top \mathbf{P}_B \mathbf{v}_j \in [0, 1]$$

- Notice also that

$$\begin{aligned} \sum_{j=1}^d \alpha_j &= \operatorname{tr}(\mathbf{V}^\top \mathbf{P}_B \mathbf{V}) = \operatorname{tr}(\mathbf{V} \mathbf{V}^\top \mathbf{P}_B) = \operatorname{tr}(\mathbf{P}_B) = \operatorname{tr}(\mathbf{B}(\mathbf{B}^\top \mathbf{B})^{-1} \mathbf{B}^\top) \\ &= \operatorname{tr}(\mathbf{B}^\top \mathbf{B}(\mathbf{B}^\top \mathbf{B})^{-1}) = \operatorname{tr}(\mathbf{I}_R) = R. \end{aligned}$$

- Thus we can consider the simplified optimization problem:

$$\text{Find } \{\alpha_j\}_{j=1}^R \text{ with } \alpha_j \in [0, 1] \text{ and } \sum_{j=1}^d \alpha_j = R \text{ that maximizes } \sum_{j=1}^d \alpha_j \lambda_j$$

# A simple proof of Eckart-Young (cont.)

- To intuitively solve the optimization problem ...

Find  $\{\alpha_j\}_{j=1}^R$  with  $\alpha_j \in [0, 1]$  and  $\sum_{j=1}^d \alpha_j = R$  that maximizes  $\sum_{j=1}^d \alpha_j \lambda_j$

- Think of  $\alpha_j$  as a purchasing variable and  $\lambda_j$  as a reward for buying the  $j$ th item
  - You must buy between 0 and 1 units of each item, and  $R$  units total
  - Question: Which purchase is the most rewarding?
  - Answer: One unit each of the  $R$  best items! i.e.,  $\alpha_j = \begin{cases} 1 & \text{if } j = 1 \dots R \\ 0 & \text{if } j = R+1 \dots d \end{cases}$
  - Recall that  $\{\lambda_j\}$  are ordered from large to small
- If  $\mathbf{v}_j$  denotes the  $j$ th eigenvector of  $\mathbf{Q}$ , these optimal  $\{\alpha_j\}$  are attained when

$$\mathbf{B} = [\mathbf{v}_1, \dots, \mathbf{v}_R] \triangleq \mathbf{V}_R, \quad \text{since } \alpha_j = \mathbf{v}_j^\top \mathbf{P}_B \mathbf{v}_j = \begin{cases} 1 & \text{if } j = 1 \dots R \\ 0 & \text{if } j = R+1 \dots d \end{cases}$$

$$\Rightarrow \mathbf{P}_B = \mathbf{V}_R \mathbf{V}_R^\top$$



# Summary of PCA



- Summary: Given centered data  $\{\mathbf{x}_i\}_{i=1}^n$ , PCA approximates  $\mathbf{x}_i$  as

$$\hat{\mathbf{x}}_i \approx \mathbf{V}_R \hat{\mathbf{z}}_i \quad \text{with} \quad \hat{\mathbf{z}}_i = \mathbf{V}_R^\top \mathbf{x}_i$$

where  $\mathbf{V}_R$  contains the  $R$  principal eigenvectors of the sample covariance mtx

$$\mathbf{Q} = \frac{1}{n} \sum_{i=1}^n \mathbf{x}_i \mathbf{x}_i^\top = \frac{1}{n} \mathbf{X}^\top \mathbf{X} \quad \text{with} \quad \mathbf{X} = [\mathbf{x}_1, \dots, \mathbf{x}_n]^\top \in \mathbb{R}^{n \times d}$$

- These eigenvectors are called the “**principal components**”
- The PCA approximation projects  $\mathbf{x}_i \in \mathbb{R}^d$  onto the subspace spanned by the  $R$  principal components of  $\mathbf{Q}$



# Performance of PCA

- How do we quantify the performance of PCA for a given rank  $R$ ?
  - This will help in choosing  $R$

- In scalar linear regression, we used  $R^2 \triangleq 1 - \frac{\text{RSS}}{ns_y^2}$

- For PCA, we will use **proportion of variance**,  $\text{PoV} \triangleq 1 - \frac{\text{RSS}}{n \text{tr}(\mathbf{Q})}$

- where  $n \text{tr}(\mathbf{Q}) = n \text{tr}(\mathbf{V} \mathbf{\Lambda} \mathbf{V}^T) = n \sum_{j=1}^d \lambda_j$

- $\text{RSS} = n \text{tr}(\mathbf{P}_B^\perp \mathbf{Q}) = n \text{tr}(\mathbf{Q}) - n \text{tr}(\mathbf{P}_B \mathbf{Q}) = n \sum_{j=1}^d \lambda_j - n \sum_{j=1}^R \lambda_j = n \sum_{j=R+1}^d \lambda_j$

- Thus the PoV for  $R$  principal components is

$$\text{PoV}(R) = \frac{n \sum_{j=1}^d \lambda_j}{n \sum_{j=1}^d \lambda_j} - \frac{n \sum_{j=R+1}^d \lambda_j}{n \sum_{j=1}^d \lambda_j} = \frac{\sum_{j=1}^R \lambda_j}{\sum_{j=1}^d \lambda_j} \quad \dots \text{want close to } 1$$

# Outline

- Dimensionality Reduction
- Principal Component Analysis (PCA)
- Computing PCA via the SVD
- Python Example: Eigenfaces and PCA-based Classification
- Data Visualization using PCA, t-SNE, and UMAP

# The singular value decomposition (SVD)

Given *any* matrix  $\mathbf{X} \in \mathbb{R}^{n \times d}$ , and denoting  $r \triangleq \text{rank}(\mathbf{X})$ :

- The standard **SVD** decomposes  $\mathbf{X}$  using square  $\mathbf{U}$  &  $\mathbf{V}$  as follows:

$$\mathbf{X} = \mathbf{U} \mathbf{S} \mathbf{V}^T \quad \text{where} \quad \begin{cases} \mathbf{U} \in \mathbb{R}^{n \times n} & \text{obeys } \mathbf{U} \mathbf{U}^T = \mathbf{I}_n = \mathbf{U}^T \mathbf{U} \\ \mathbf{V} \in \mathbb{R}^{d \times d} & \text{obeys } \mathbf{V} \mathbf{V}^T = \mathbf{I}_d = \mathbf{V}^T \mathbf{V} \\ \mathbf{S} \in \mathbb{R}^{n \times d} & \text{obeys } \mathbf{S} = \text{Diag}(s_1, \dots, s_m) \\ & \text{where } m = \min\{n, d\} \\ & \text{and } s_1 \geq s_2 \geq s_3 \geq \dots \geq 0 \end{cases}$$

- The “**economy SVD**” decomposes  $\mathbf{X}$  using square  $\mathbf{S}_r$  and tall  $\mathbf{U}_r$  &  $\mathbf{V}_r$ :

$$\mathbf{X} = \mathbf{U}_r \mathbf{S}_r \mathbf{V}_r^T \quad \text{where} \quad \begin{cases} \mathbf{U}_r \in \mathbb{R}^{n \times r} & \text{obeys } \mathbf{U}_r^T \mathbf{U}_r = \mathbf{I}_r \\ \mathbf{V}_r \in \mathbb{R}^{d \times r} & \text{obeys } \mathbf{V}_r^T \mathbf{V}_r = \mathbf{I}_r \\ \mathbf{S}_r \in \mathbb{R}^{r \times r} & \text{obeys } \mathbf{S}_r = \text{Diag}(s_1, \dots, s_r) \\ & \text{where } r \leq \min\{n, d\} \\ & \text{and } s_1 \geq s_2 \geq \dots \geq s_r > 0 \end{cases}$$

# Computing PCA via the standard SVD

- As before, assume that  $\mathbf{X} = [\mathbf{x}_1, \dots, \mathbf{x}_n]^\top \in \mathbb{R}^{n \times d}$  is centered
  - That is, the sample mean of every column equals zero
- For PCA, we used the (sorted) eigen-decomposition of the covariance matrix

$$\mathbf{Q} = \frac{1}{n} \mathbf{X}^\top \mathbf{X} = \mathbf{V} \mathbf{\Lambda} \mathbf{V}^\top \quad \text{where} \quad \begin{cases} \mathbf{V} \mathbf{V}^\top = \mathbf{I}_d = \mathbf{V}^\top \mathbf{V} \\ \mathbf{\Lambda} = \text{Diag}(\lambda_1, \dots, \lambda_d) \\ \lambda_1 \geq \lambda_2 \geq \lambda_3 \geq \dots \geq 0 \end{cases}$$

- Plugging in the standard SVD of  $\mathbf{X}$ , i.e.,  $\mathbf{X} = \mathbf{U} \mathbf{S} \mathbf{V}^\top$ , we find

$$\mathbf{Q} = \frac{1}{n} \mathbf{V} \mathbf{S}^\top \underbrace{\mathbf{U}^\top \mathbf{U}}_{\mathbf{I}_n} \mathbf{S} \mathbf{V}^\top = \mathbf{V} \left( \frac{1}{n} \mathbf{S}^\top \mathbf{S} \right) \mathbf{V}^\top \quad \text{where} \quad \begin{cases} \mathbf{V} \mathbf{V}^\top = \mathbf{I}_d = \mathbf{V}^\top \mathbf{V} \\ \frac{1}{n} \mathbf{S}^\top \mathbf{S} = \text{Diag}\left(\frac{s_1^2}{n}, \dots, \frac{s_d^2}{n}\right) \\ \frac{s_1^2}{n} \geq \frac{s_2^2}{n} \geq \frac{s_3^2}{n} \geq \dots \geq 0 \end{cases}$$

- So, the quantities  $\mathbf{V}, \{s_j\}$  from the SVD of  $\mathbf{X}$  are sufficient to compute PCA:
  - $\lambda_j = s_j^2/n$  and the  $\mathbf{V}$  matrices are the same (if eigenvals are distinct and sorted)

# Computing PCA via the economy SVD

- Remember: the economy SVD computes only the top  $r$  singular vectors, i.e.,  $U_r$  and  $V_r$ , where  $r = \text{rank}(\mathbf{X}) \leq \min(n, d)$ 
  - Thus, it may require less computational complexity than the standard SVD
- For PCA, need to compute only the top  $R$  singular vectors  $V_R$ , where  $R < r$ 
  - $R$  is a design choice, and typically  $R \ll r$
- Thus it's more efficient to use the economy SVD when computing PCA:

$$(U_r, S_r, V_r^T) = \text{SVD}_{\text{economy}}(\mathbf{X})$$

$$V_R = \text{first } R \text{ columns of } V_r$$

$$\mathbf{z}_i = V_R^T \mathbf{x}_i \quad \forall i = 1 \dots n$$

# Standardizing the PCA coefficients

- After computing the PCA coefficients  $\{z_i\}$ , we might use them for classification or regression (assuming that we also have some targets  $\{y_i\}$ )
- If so, we should **standardize** our new features  $\{z_i\}$ . For this, we want
  - $\bar{z} \triangleq \frac{1}{n} \sum_{i=1}^n z_i = 0$
  - $\frac{1}{n} \sum_{i=1}^n z_{ij}^2 = 1$  for all  $j = 1 \dots R \quad \Leftrightarrow \quad \text{Diag} \left( \overbrace{\frac{1}{n} \sum_{i=1}^n z_i z_i^T}^{\triangleq Q_z} \right) = 1$
- Let's analyze the actual values of  $\bar{z}$  and  $\text{Diag}(Q_z)$ :
  - $\bar{z} = \frac{1}{n} \sum_{i=1}^n V_R^T x_i = V_R^T (\frac{1}{n} \sum_{i=1}^n x_i) = V_R^T \mathbf{0} = \mathbf{0}$ , since  $\{x_i\}$  are centered
  - $Q_z = \frac{1}{n} \sum_{i=1}^n z_i z_i^T = V_R^T (\frac{1}{n} \sum_{i=1}^n x_i x_i^T) V_R = V_R^T Q V_R = \frac{1}{n} V_R^T V_r S_r^2 V_r^T V_R = \text{Diag}(\frac{s_1^2}{n}, \dots, \frac{s_R^2}{n})$  since  $V_R^T V_r = [I_R \quad \mathbf{0}_{R \times (r-R)}]$
- This implies that a standardized version of  $\{z_i\}$  is given by  $\{\tilde{z}_i\}$  with
  - $\tilde{z}_i \triangleq \text{Diag} \left( \frac{\sqrt{n}}{s_1}, \dots, \frac{\sqrt{n}}{s_R} \right) z_i$

# Outline

- Dimensionality Reduction
- Principal Component Analysis (PCA)
- Computing PCA via the SVD
- Python Example: Eigenfaces and PCA-based Classification
- Data Visualization using PCA, t-SNE, and UMAP



# PCA on the LFW face dataset

- First we center the data  $\{x_i\}$

```
Xmean = np.mean(X,0)
Xs = X - Xmean[None,:]
```

- Then we compute the economy SVD


```
U,S,VT = np.linalg.svd(Xs, full_matrices=False)
VT.shape

(1288, 1850)
```

- Note that  $V_r^T$  is wide, as we expect


```
lam = S**2 / n_samples
PoV = np.cumsum(lam)/np.sum(lam)
```

- Then we compute the eigenvalues  $\{\lambda_j\}_{j=1}^r$
- And finally we compute the Proportion-of-Variance
- The PoV plot suggests that  $R = 400$  principal components capture nearly all the variance of our data




# PCA approximation and eigenfaces

- We now show the PCA approximations versus  $R$  for two faces:



- And the mean & top 5 principal components  $\{v_j\}$  (i.e., “eigenfaces”):



# Face recognition using the PCA coefficients

We now demonstrate classification (i.e., face recognition) via PCA coefficients

- Split data into training and test:

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25, stratify = y, random_state=43)
```

- Center the training:

- Perform economy SVD:

```
n_samples, _ = X_train.shape
Xtr_mean = np.mean(X_train,0)
Xtr = X_train - Xtr_mean[None,:]
Utr,Str,VTtr = np.linalg.svd(Xtr, full_matrices=False)
```

- Choose  $R = 100$ :

- Compute PCA coefficients

$$\mathbf{z}_i^T = \mathbf{x}_i^T \mathbf{V}_R \quad \forall i:$$

```
npc = 100
eigenfaces = VTtr[:npc,:]
Ztr = Xtr.dot(eigenfaces.T)
```

- Standardize the PCA

$$\text{coefficients: } \tilde{\mathbf{z}}_i = \text{Diag}\left(\frac{\sqrt{n}}{s_1}, \dots, \frac{\sqrt{n}}{s_R}\right) \mathbf{z}_i:$$

```
Ztr_s = Ztr / Str[None,:npc] * np.sqrt(n_samples)
```

## Face recognition using the PCA coefficients (cont.)

- Tune an SVM classifier over regularization  $C$  & RBF kernel width  $\gamma$ :

```
param_grid = {'C': [1, 3, 10, 30, 100, 300],
              'gamma': [0.00001, 0.00003, 0.0001, 0.0003, 0.001, 0.003, 0.01], }
clf = GridSearchCV(SVC(kernel='rbf', class_weight='balanced'), param_grid, cv=5, iid=False)
clf = clf.fit(Ztr_s, y_train)
print("Best estimator found by grid search:")
print(clf.best_estimator_)
print("Cross validation accuracy with best estimator:")
print(clf.best_score_)
```

```
Best estimator found by grid search:
SVC(C=3, cache_size=200, class_weight='balanced', coef0=0.0,
    decision_function_shape='ovr', degree=3, gamma=0.003, kernel='rbf',
    max_iter=-1, probability=False, random_state=None, shrinking=True,
    tol=0.001, verbose=False)
Cross validation accuracy with best estimator:
0.8506081737123565
```


- After pre-processing the test data in the same way as training, classify it:

```
Xts = X_test - Xtr_mean[None,:]
Zts = Xts.dot(eigenfaces.T)
Zts_s = Zts / Str[None,:npc] * np.sqrt(n_samples)
y_hat = clf.predict(Zts_s)
acc = np.mean(y_hat==y_test)
print("The model accuracy on the test set is %f" % acc)
```

```
The model accuracy on the test set is 0.829193
```

# Limitations of PCA

- As a dimensionality reduction technique, PCA has two main characteristics:  
1) it is linear, and 2) it minimizes RSS.
  - These characteristics are convenient and allow us to derive a closed-form solution for PCA
- But minimizing RSS is not well justified for tasks like classification
  - For example, PCA can destroy the linear separability of a dataset, as illustrated here:
- **Linear discriminant analysis (LDA)** is a different form of linear dimensionality reduction that explicitly aims to discriminate between two classes
  - Although it's a nice idea, LDA isn't widely used today

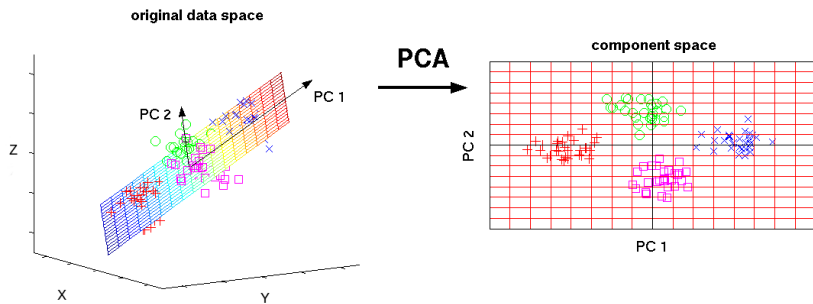


# Outline

- Dimensionality Reduction
- Principal Component Analysis (PCA)
- Computing PCA via the SVD
- Python Example: Eigenfaces and PCA-based Classification
- Data Visualization using PCA, t-SNE, and UMAP


# Data visualization using PCA

- When  $d \geq 3$ , the scatter plot of  $\{x_i\}_{i=1}^n$  can be difficult to visualize
- But when  $R \leq 2$ , the scatter plot of  $\{z_i\}_{i=1}^n$  can be easily visualized:



- The principal components  $\{v_j\}_{j=1}^R$  may also be meaningful ...


# Example: PCA visualization of MNIST





# Data visualization via t-SNE and UMAP

- To **visualize** high-dimensional data, it is best to use a **nonlinear dimensionality reduction** approach
- Classical methods include **kernel PCA**, **IsoMap**, and **locally linear embedding**, but they don't work very well
- The most popular methods today include **t-stochastic neighbor embedding (t-SNE)**, **uniform manifold approximation and projection (UMAP)**, and **PaCMAP**



# Learning objectives

- Recognize need for feature dimensionality reduction
- Understand PCA as RSS-minimizing linear approximation
  - Understand orthogonal projection
  - Recognize PCA as subspace fitting
  - Understand the role of the data-covariance eigenvectors in PCA
  - Know how to measure PCA performance using PoV
- Understand how to compute PCA using the SVD
- Understand how the PCA coefficients can be used in supervised learning tasks
- Understand how PCA can be used for data visualization
- Be familiar with t-SNE and UNET, non-linear data-visualization techniques